

Page 1: Introduction to RAG (Retrieval Augmented Generation)

RAG is a technique that combines information retrieval and text generation using LLMs.

It retrieves relevant documents from a vector database and sends them to an LLM.

Page 2: Components of RAG

1. Loader: Loads documents like PDF.
2. Text Splitter: Splits documents into chunks.
3. Embeddings: Converts text into vectors.
4. FAISS: Stores and searches vectors efficiently.

Page 3: Working Flow

PDF → Chunking → Embedding → FAISS Storage → Retrieval → LLM Answer

This pipeline is used in modern AI chatbots and search systems.